Application of HPX to Tiled GEMM and QR: A Benchmark

September 25, 2019 | Thomas Miethlinger | Jülich Supercomputing Centre

JÜLICH
Forschungszentrum

Part I: Introduction

JÜLICH
Forschungszentrum

About me

(Thomas Miethlinger)

- Study: Master Physics
- Johannes Kepler University of Linz
- Institute for Theoretical Physics
  Department: Many Particle Systems
  Research:
  - Quantum fluids
  - Complex fluids
  - Non-equilibrium statistical mechanics
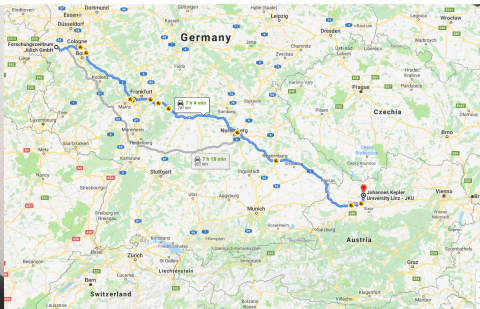
About the Supervisors

- Supervisor: Dr. Edoardo Di Napoli
- Co-Supervisor: Dr. Xinzhe Wu
- SimLab Quantum Materials
- Research:
    - Development and maintenance of numerical libraries
    - Design and implementation of high-performance algorithms
    - Development of new mathematical and computational models within a methodological framework

  in the scope of computational materials science and quantum materials.

JÜLICH
Forschungszentrum
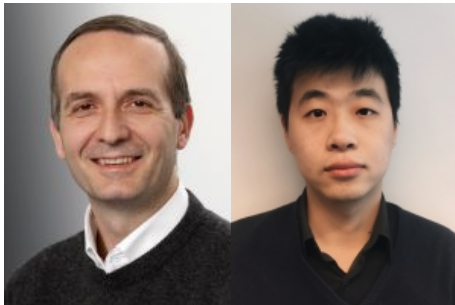
Part II: Introduction to HPX

JÜLICH
Forschungszentrum

Current situation in high performance computing (HPC)

Currently, speed-up in computing does not stem from higher CPU frequency, but increased parallelism.
However, we already face the following challenges in HPC:

- Ease of programming
- Inability to handle dynamically changing workloads
- Scalability
- Efficient utilization of system resources
- $\Longrightarrow$ a need for a new execution model: ParalleX, which is implemented by HPX

JÜLICH
Forschungszentrum

ParalleX

ParalleX is a new parallel execution model that offers an alternative to the conventional computation models(e.g. message passing):

- Split-phase transaction model
- Message-driven
- Distributed shared memory
- Multi-threaded
- Futures synchronization
- Local Control Objects (LCOs)
- ...

ParalleX focuses on latency hiding instead of latency avoidance.

JÜLICH
Forschungszentrum

About HPX

- High Performance ParalleX (HPX) is the first runtime system implementation of the ParalleX execution model.
- Development: STE||AR group
  Louisiana State University
  LSU Center for Computation and Technology
- Released as open source under the Boost Software License
- Current version: HPX V1.3.0, released on 23.05.2019
- Aims to be a **C++ standards conforming implementation** of the Parallelism and Concurrency proposals for C++ 17/20/23/...
- This means: HPX is a C++ library that supports **dynamic adaptive resource management** and **lightweight task programming and scheduling** within the context of a **global address space**.

JÜLICH
Forschungszentrum

On learning HPX

An opinion of a non-CS/HPC student

Learning curve on of HPX is quite steep - in the first days quite some dedication, effort and endurance is needed[1].

- Probably the easiest way in the beginning: watch this nice playlist in 1.25x speed on the Youtube channel of cscsch (Swiss National Supercomputing Centre)
- Be aware that the API reference is not complete
- Be aware that there exist at least 5 different "Hello, World!" examples[2]:
    - hpx/examples/hello_world_component/*: 3 files; 28, 30 & 55 lines
    - hpx/examples/quickstart/hello_world_1.cpp; 22 lines
    - hpx/examples/quickstart/hello_world_2.cpp; 24 lines
    - hpx/examples/quickstart/hello_world_distributed.cpp; 156 lines
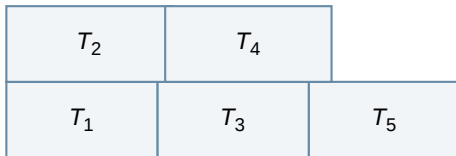   - tutorials/examples/01_hello_world/hello_world.cpp; 71 lines

---

[1]Why is the HPX code repo so big and complicated?
[2]Paths are with respect to https://github.com/STEllAR-GROUP/

JÜLICH
Forschungszentrum

HPX: Tasks and Threads

- HPX: Task-based parallelism
- Split up big problem into smaller tasks
- Tasks are worked off as HPX (lightweight) Threads by the OS Threads
- Task size is crucial: not too small and not too big
- Number of tasks can even be as high as $\mathcal{O}(10^8)$

Tasks too large

| T_2	T_4	
T_1	T_3	T_5

Right task size

T_2	T_4	T_5	T_9	
T_1	T_3	T_6	T_7	T_8

JÜLICH
Forschungszentrum
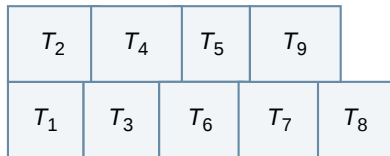
Comparison of HPX and OpenMP

HPX	OpenMP
C++ library	Compiler extension to C and Fortran
Core language: `hpx::C++`	`#pragma omp` directives
Task-based parallelism	Parallel regions (fork-join model)

JÜLICH
Forschungszentrum

Comparison of HPX and MPI

HPX	MPI
C++ library	Interface specification for C and Fortran
Core language: `hpx::C++`	Core language: MPI_C, MPI_F08
Task-based parallelism	Single program, multiple data (SPMD)
Message driven	Message passing

JÜLICH
Forschungszentrum

The HPX API

Selection: Classes

Class	Description
`hpx::thread`	Low level thread of control
`hpx::mutex`	Low level synchronization facility
`hpx::lcos::local::condition_variable`	Signal a condition
`hpx::future, hpx::shared_future`	Asynchronous result transport (receiving end)
`hpx::promise, hpx::lcos::local::promise`	Asynchronous result transport (producing end)
`hpx::lcos::packaged_task`	Asynchronous result transport (producing end)
`hpx::function`	Type erased function object
`hpx::tuple`	Tuple
`...`	

JÜLICH
Forschungszentrum

The HPX API

Selection: Functions

Functions	Description
`hpx::async`	Spawning tasks (returns a future)
`hpx::make_ready_future`	Spawning tasks (returns a ready future)
`hpx::bind`	Binding Parameters to callables
`hpx::apply`	Signal a condition
`future::{is_ready, valid, has_exception}`	Query state of future
`future::get`	Return computed result of future
`future::then`	Continuations of futures
`hpx::when_all, hpx::when_any, hpx::when_n`	Waiting on one or more futures (non blocking)
`hpx::wait_all, hpx::wait_any, hpx::wait_n`	Waiting on one or more futures (blocking)
`hpx::dataflow`	Shortcut to hpx::when_all(...).then(...)
`...`	

JÜLICH
Forschungszentrum

HPX: Example Program

```cpp
double calc_area(hpx::future<double> future_r, hpx::future<double> future_pi)
{
    double r = future_r.get(); // r is returned immediately (make_ready_future)
    double pi = future_pi.get(); // pi is returned once the async computation finishes
    return r * r * pi; // return area = r^2 * pi
}

int hpx_main(variables_map& vm) // In hpx_main the HPX environment is loaded
{   // boost::program_options::variables_map: retrieve commandline arguments
    hpx::future<double> future_r = hpx::make_ready_future(vm["r"].as<double>());
    hpx::future<double> future_pi = hpx::async([](){ return 4.0 * atan(1.0); });
    hpx::future<double> future_area = hpx::dataflow(&calc_area, future_r, future_pi);
    return hpx::finalize(); // Area can be obtained by: future_area.get()
}

int main(int argc, char * argv[]) // Start program by: ./area --r=...
{   // boost::program_options::options_description: handle commandline arguments
    options_description.add_options()("r", value<double>()->default_value(1.0), "Radius: r");
    return hpx::init(options_description, argc, argv); // hpx::init calls hpx_main
}
```

JÜLICH
Forschungszentrum

Part III: Introduction to Numerical Linear Algebra, GEMM and QR

JÜLICH
Forschungszentrum

Introduction: Numerical Linear Algebra

Numerical linear algebra is a subfield of numerical analysis and linear algebra, and it plays an integral role in computational problem solving. There exist many several algorithms for common problems, a few well-known are:

- Solving systems of linear equations
- Eigenvalue problem
- Matrix inversion problem
- Least-squares problem

which may be using one of the following matrix operations/decompositions:

- Matrix multiplications
- LU decomposition
- QR decomposition
- Spectral decomposition
- Singular value decomposition

GEMM

Applications of Numerical Linear Algebra

GEMM - GEneral Matrix Multiply

- Basic binary operation in Linear Algebra, which has numerous applications in mathematics, science and engineering.
- More fundamental applications of matrix multiplications include
  - Systems of Linear Algebraic Equations (SLAE) can be expressed as a single matrix equation, e.g. $Ax = y$.
  - Linear map between two vector spaces U and V over the same field F.
- Motivation: A large amount (>70%) of runtime in ChASE ((E.D. Napoli, 2019)) routine is used for GEMM.
- Let the field F be $\mathbb{R}$ or $\mathbb{C}$, $A = (a_{ij}) \in F^{m \times n}$, $B = (b_{jk}) \in F^{n \times p}$. Then,

$$C = (c_{ik}) = AB \in F^{m \times p}, \tag{1}$$

$$c_{ik} = \sum_{j=1}^{n} a_{ij} b_{jk} \tag{2}$$

JÜLICH
Forschungszentrum

GEMM of Blocked Matrices

Applications of Numerical Linear Algebra

$\implies$ most simple implementation consists of 3 nested for-loops:

```
for 0 <= i < m, 0 <= j < n, 0 <= k < p, do: C[i][k] += A[i][j] * B[j][k]
```

Better approach: Discretize matrices into blocks, perform GEMM **block-wise**

Let the field F again be $\mathbb{R}$ or $\mathbb{C}$, $A = (A_{ij}) \in F^{M \times N \times m \times n}$, $B = (B_{jk}) \in F^{N \times P \times n \times p}$.

Then:

$$C = (C_{ik}) = AB \in F^{M \times P \times m \times p}, \tag{3}$$

$$C_{ik} = (c_{ik,i'k'}) = \sum_{j=1}^{n} A_{ij} B_{jk}, \tag{4}$$

$$c_{ik,i'k'} = \sum_{j=1}^{N} \sum_{j'=1}^{n} a_{ij,i'j'} b_{jk,j'k'} \tag{5}$$

JÜLICH
Forschungszentrum

GEMM of Blocked Matrices

A Small Example

Let $m = n = p = M = N = P = 2$:

$$A \cdot B = C$$

$$
\begin{array}{|c|c|}
\hline
A_{00} & A_{10} \\
\hline
A_{01} & A_{11} \\
\hline
\end{array}
\cdot
\begin{array}{|c|c|}
\hline
B_{00} & B_{10} \\
\hline
B_{01} & B_{11} \\
\hline
\end{array}
=
\begin{array}{|c|c|}
\hline
C_{00} & C_{10} \\
\hline
C_{01} & C_{11} \\
\hline
\end{array}
$$

JÜLICH
Forschungszentrum

GEMM of Blocked Matrices

A Small Example

Let $m = n = p = M = N = P = 2$:

$$A_{ij} \cdot B_{jk} = A_{ij}B_{jk}$$

$a_{ij,00}$	$a_{ij,10}$
$a_{ij,01}$	$a_{ij,11}$

$\cdot$

$b_{jk,00}$	$b_{jk,10}$
$b_{jk,01}$	$b_{jk,11}$

$=$

...	...
...	...

JÜLICH
Forschungszentrum

GEMM of Blocked Matrices

OpenMP and HPX foreach Row

Let $M = N = 16$, number of threads $\mathcal{T} = 4$:


OMP Row Tiled & HPX foreach Row

OMP Row Cell

JÜLICH
Forschungszentrum

GEMM of Blocked Matrices

HPX future C_{ik} and HPX future $A_{ij}B_{jk}$

HPX future $f_{ik} \rightarrow C_{ik}$

...	$f_{ik} \rightarrow C_{ik}$	...

$$C_{ik} = f_{ik}.\texttt{get()}$$

HPX future $f_{ijk} \rightarrow A_{ij}B_{jk}$

...	$f_{ijk} \rightarrow A_{ij}B_{jk}$	...

$$C_{ik} = \sum_j f_{ijk}.\texttt{get()}$$

JÜLICH
Forschungszentrum

GEMM of Blocked Matrices

Distributed Memory

Let $M = N = 2$, number of nodes $\mathcal{N} = 4$:

MPI V1 & HPX dist V1

| $\mathcal{N}_0$ A_{00}, B $A_{00}B_{00}$ C_{00} | $\mathcal{N}_1$ A_{01}, B $A_{01}B_{10}$ |
| $\mathcal{N}_2$ A_{10}, B $A_{10}B_{00}$ C_{10} | $\mathcal{N}_3$ A_{11}, B $A_{11}B_{10}$ |

MPI V2 & HPX dist V2

| $\mathcal{N}_0$ A_{00}, B_{0k} $A_{00}B_{0k}$ C_{00} | $\mathcal{N}_1$ A_{01}, B_{1k} $A_{01}B_{1k}$ C_{01} |
| $\mathcal{N}_2$ A_{10}, B_{0k} $A_{10}B_{0k}$ C_{10} | $\mathcal{N}_3$ A_{11}, B_{1k} $A_{11}B_{1k}$ C_{11} |

JÜLICH
Forschungszentrum

QR decomposition

Applications of Numerical Linear Algebra

QR decomposition

- Matrix decomposition of square or rectangular matrices (we only consider square matrices).
- Let the field F be $\mathbb{R}$ or $\mathbb{C}$, $A = (a_{ij}) \in F^{m \times m}$. Then,

$$A = QR, \tag{6}$$

  where Q is a orthogonal($F = \mathbb{R}$) / unitary($F = \mathbb{C}$) matrix, and R is an upper triangular matrix.

- Motivation: A large amount (>20%) of runtime in ChASE ((E.D. Napoli, 2019)) routine is used for QR.
- Computing the QR decomposition:
  - Gram-Schmidt process
  - **Householder reflection**
  - Givens rotations

JÜLICH Forschungszentrum

Applications of QR decomposition

Examples

- Computing Eigenvalues (QR algorithm)
- Computing orthogonal base
- Solving least-squares problem
- Solution to linear inverse problems

A selection of papers within the scope of my research interest which use QR decomposition:

- An evaluation of noise reduction algorithms for particle-based fluid simulations in multi-scale applications (M.J. Zimon et al., 2016)
- Dynamic mode decomposition of numerical and experimental data (P.J. Schmid, 2010)
- Krylov Methods for the Incompressible Navier-Stokes Equations (W.S. Edwards, 1994)
- Computing Lyapunov exponents of continuous dynamical systems: method of Lyapunov vectors (J. Lu, 2005)

JÜLICH
Forschungszentrum

Block QR decomposition

Example: 3×3 **block matrix**

Block QR decomposition

Example: 3×3 **block matrix**

geqrt

tpqrt

gemqrt

tpmqrt

JÜLICH
Forschungszentrum
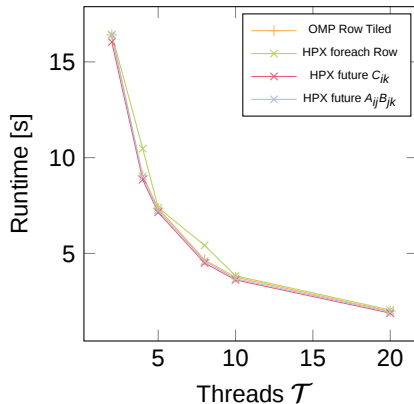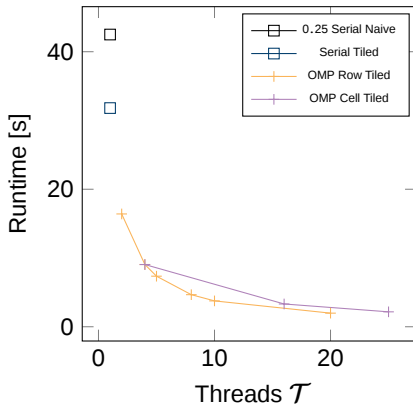
Part IV: Benchmark: Results for GEMM

GEMM Shared Memory

Let $m = n = p = 4000$, $M = N = P = 40$, vary number of threads $\mathcal{T}$:

GEMM Shared Memory

Let $m = n = p = 4000$, $M = N = P = 40$, vary number of threads $\mathcal{T}$:

$\implies$ HPX can be faster than OpenMP, even for embarrassingly parallel problems, in particular `hpx::futures`.

JÜLICH
Forschungszentrum

GEMM Shared Memory

Vary number of threads $\mathcal{T}$, and let $m = n = p = 400\sqrt{\mathcal{T}}, M = N = P = \sqrt{\mathcal{T}}$:


$\implies$ Task-size crucial for HPX. Here the block size was too big.
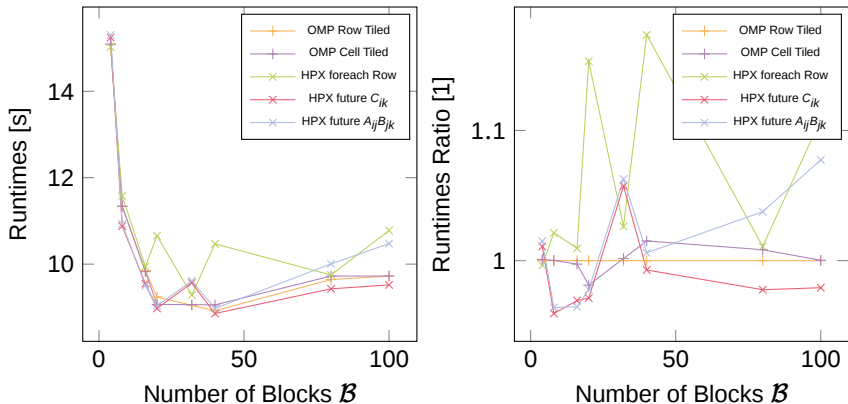
GEMM Distributed Memory

Vary number of Nodes $\mathcal{N}$, and let $m = n = p = 400\sqrt{\mathcal{N}}, M = N = P = \sqrt{\mathcal{N}}$:

HPX distributed is slower than MPI, their implementations are presumably analogous.

JÜLICH
Forschungszentrum

GEMM Shared Memory

Let $m = n = p = 4000$, $\mathcal{T} = 4$, vary number of blocks $\mathcal{B} := M = N = P$:

$\implies$ Again, HPX is around the same speed as OpenMP.

Part V: Benchmark: Results for Tiled QR

QR Shared Memory

Let $m = 4000$, $M = 40$, vary number of threads $\mathcal{T}$:

QR Shared Memory

Vary number of threads $\mathcal{T}$, and let $m = n = p = 400\sqrt{\mathcal{T}}, M = N = P = \sqrt{\mathcal{T}}$:

QR Shared Memory

Let $m = n = p := 4000$, $\mathcal{T} = 4$, vary number of blocks $\mathcal{B} := M = N = P$:

$\Longrightarrow$ Discretizing problem into blocks has a high impact on runtime!

Part VI: Conclusion and Resume

Conclusion
On HPX

- HPX has the potential to speed-up and simplify parallelization for certain types of problems
- On par with OpenMP
- Comparison with MPI unfortunately not possible at this moment (supplemented for report)
- Learning HPX takes some time, in particular HPX for distributed programs
- Debugging, profiling, runtime errors, documentation, ... could be improved
- Execution policies: might give additional speed-up
- Altogether: Very strong library and concept with some negative points

JÜLICH
Forschungszentrum

Resume

Personal thoughts

- Interesting project, work and topic. I learned a lot.
- Collaboration with Dr. Wu worked very well. Thanks!
- On HPX
    - * Some fun
    - * Takes a medium amount of time
    - \+ Full C++ library
    - – Bugs can be nasty
- On MPI
    - \+ Lots of fun
    - – Takes a lot of time
    - – No native C++
    - * Usually works, and if not, it is manageable to figure out why
- On OpenMP
    - – No fun
    - \+ Doesn't take a lot of time
    - \+ Simply works :)

JÜLICH
Forschungszentrum

Application of HPX to Tiled GEMM and QR: A Benchmark

September 25, 2019 | Thomas Miethlinger | Jülich Supercomputing Centre

JÜLICH
Forschungszentrum